

FraudShield

Real-Time Credit Card Fraud Detection Gateway

Project Synopsis

Submitted by: Ayush Sharma

Problem Statement

Rule-based fraud detection — fixed thresholds on amount, velocity, or location — catches only what it was explicitly written to catch. Tuned loosely, it misses novel fraud patterns; tuned aggressively, it blocks large numbers of legitimate transactions and creates real customer friction. Most small and mid-sized platforms lack a system that combines the speed of simple rules with a trained model that generalizes beyond them, routes genuinely ambiguous cases to a human analyst instead of guessing, and treats the detection threshold as a deliberate, documented trade-off rather than a hardcoded default.

Proposed System

FraudShield is a full-stack fraud-detection gateway serving four roles — Customer, Merchant, Fraud Analyst, and Admin. Each simulated transaction passes through a rule-based pre-filter, then a trained classifier that returns a fraud-probability score; the combined result routes it to auto-approve, auto-decline, or a Fraud Analyst's review queue, where every decision is logged as labeled data for future retraining. Admins can adjust the decision threshold directly, with the underlying trade-off made explicit rather than hidden. The classifier is trained separately on two datasets serving different purposes — IEEE-CIS Fraud Detection (590,540 real-world e-commerce transactions) as the primary production model, and a large, Sparkov-simulated transactions dataset ($\approx 1.85\text{M}$ rows) with fully readable transaction-category, merchant, and location columns, kept specifically for SHAP-based explainability.

Objectives

- Train and evaluate a machine-learning fraud classifier on real-world-scale transaction data, using metrics suited to severe class imbalance — precision, recall, F1-score, and AUC-ROC — rather than raw accuracy.
- Combine a fast rule-based pre-filter with the trained model into a single real-time scoring pipeline.
- Build a role-based web application so Customers, Merchants, Fraud Analysts, and Admins each get the view and permissions appropriate to their role.
- Give Fraud Analysts a prioritized review queue, and Admins a way to adjust — and justify — the detection threshold.
- Track model performance over time as analyst-labeled outcomes accumulate.
- Deploy a working, end-to-end version of the system for live demonstration.

Scope

In scope: a simulated transaction gateway (not a live payment processor), the full rule-based + ML detection pipeline, all four role-based interfaces, threshold management, performance analytics, and analyst notifications. Model training draws on the IEEE-CIS dataset as the primary source and a Sparkov-simulated transactions dataset as a secondary, explainability-focused source, trained separately rather than merged since the two feature schemas are incompatible.

Out of scope for this build, left as future enhancements: real (sandboxed or live) payment-gateway integration, streaming architectures such as Kafka, deep-learning approaches such as autoencoders or graph neural networks, and continuous/online learning.

Technology Stack

- Frontend: React (Vite), Tailwind CSS
- Backend: Django, Django REST Framework
- Authentication: Django auth + djangorestframework-simplejwt (JWT), role-based permissions
- ML / Data: scikit-learn, pandas, NumPy, imbalanced-learn (SMOTE), joblib
- Database: PostgreSQL, via Neon (pooled connection) in production
- Visualization: Recharts
- Notifications: Twilio / Django email
- Deployment: Vercel (frontend + backend) + Neon (database)

Expected Outcomes

- A deployed web application that scores simulated transactions in real time and routes them to approve, decline, or analyst review.
- A trained classifier that measurably outperforms the rule-based pre-filter alone on held-out test data, reported via precision, recall, F1-score, and AUC-ROC.
- A working Fraud Analyst workflow where review decisions are captured as labeled data for future retraining.
- An Admin-configurable detection threshold with a documented rationale for the value chosen.
- A performance-analytics view showing model metrics trending as labeled data accumulates.
- A live demo in which the trained model catches at least one fraud case a rule-only check would have missed.

Timeline

Week	Focus
1	Environment & project scaffolding
2	Data models & migrations
3	Authentication & role-based access
4	Dataset acquisition & exploratory data analysis
5	Model training pipeline
6	Model serving & rule-based pre-filter integration
7	Transaction simulation module
8	Merchant & Fraud Analyst dashboards
9	Admin panel & threshold management
10	Model performance analytics
11	Notifications
12	Testing, deployment & demo prep